NVIDIA

Toward Real-Time Generation With MiniMax H3

8× RTX PRO 5000 · vLLM-Omni

Shunyang Li | Benchmarks measured | 27 Oct 2026



Agenda

01 Toward Real-Time H3 Generation

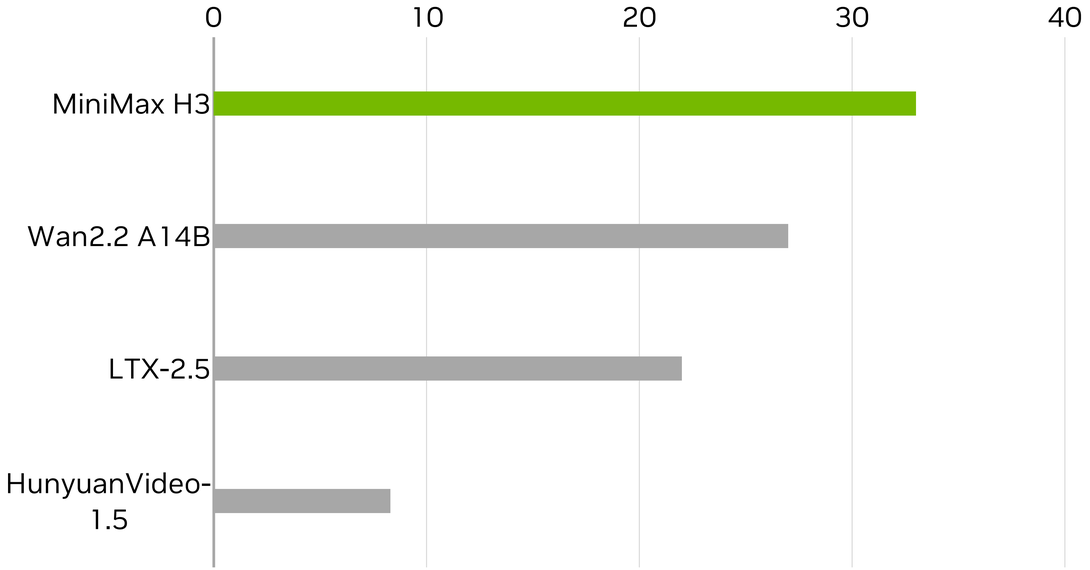
02 vLLM-Omni and ComfyUI Applications



01 · Real-Time H3 Generation

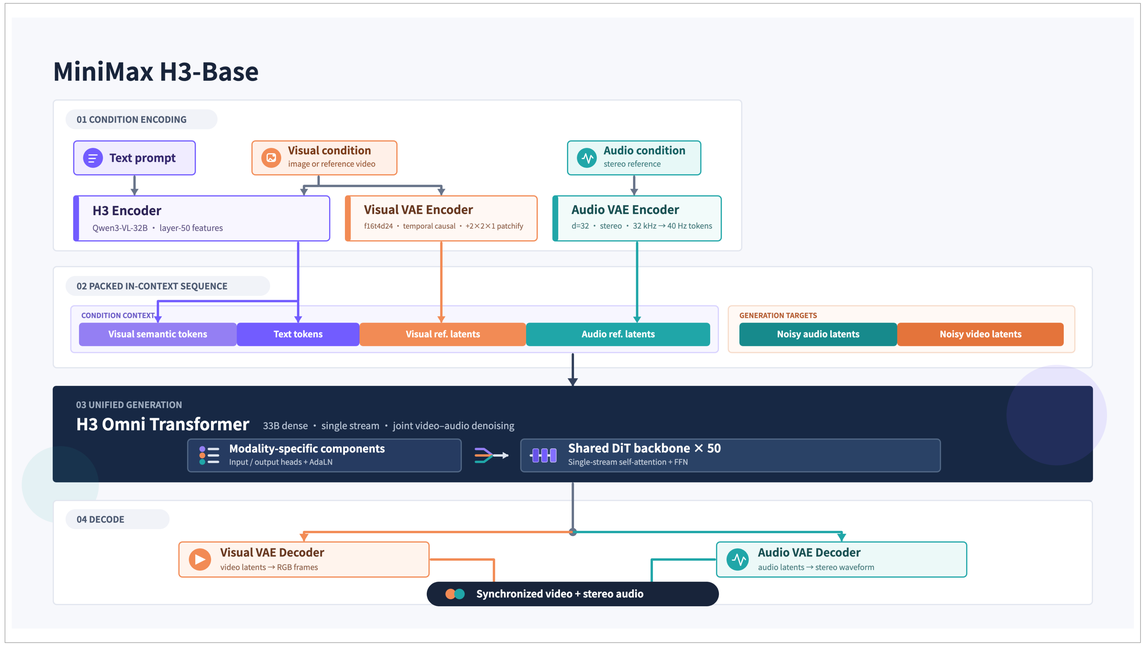
Real-Time Ref2VA Raises the Systems Bar

A 124-frame clip at 24 FPS is real time only when end-to-end latency stays below 5.167 s (RTF < 1)



Generative Backbone · Full Parameters (B)

RTF = latency ÷ media duration; the 124-frame deadline is 5.167 s. H3 has the largest dense backbone in this peer set; Wan2.2 has 27B total and 14B active parameters



Model Capability · Native 768P · 33B Dense

A 32B encoder feeds the 33B Omni Transformer and visual/audio VAEs. Fifty transformer blocks are evaluated at each of 50 baseline steps.

Sources: MiniMax H3, Wan2.2, LTX-2.5 and HunyuanVideo-1.5 official releases · author 960×544 Ref2VA workload

NVIDIA

SM120 Brings FP4 Compute; PCIe Limits Scale

All discrete SM120 cards pair fifth-generation Tensor Cores with GDDR7 and PCIe Gen5; none provides GPU-to-GPU NVLink

SKU	Compute	Memory	GPU I/O / Fabric
RTX PRO 5000 ×8 Blackwell	SM120 · 110 SMs 5th-gen Tensor · FP4	72 GB GDDR7 ECC 48 GB variant	PCIe Gen5 ×16 Planned: RDMA via NIC*
RTX PRO 6000 WS / Server	SM120 · 188 SMs 5th-gen Tensor · FP4	96 GB GDDR7 ECC	PCIe Gen5 ×16 No GPU NVLink
GeForce RTX 5090	SM120 · 170 SMs 5th-gen Tensor · FP4	32 GB GDDR7	PCIe Gen5 ×16 No NVLink
DGX Spark · GB10	SM121 · 48 SMs 5th-gen Tensor · FP4	128 GB unified LPDDR5X	NVLink-C2C in SoC 200 Gb/s ConnectX-7

* Planned system path; GPUDirect RDMA requires a compatible NIC and PCIe topology · Sources: NVIDIA product specs, CUDA GPU table, GPUDirect RDMA guide

NVIDIA

Baseline: Nsight Systems Breakdown

Original Ref2VA · 50 evaluations · BF16 · dense attention · 8 × SM120 GPUs

BASELINE E2E

213.402 s

Complete-MP4 HTTP median · n = 2

PROFILED REQUEST

214.968 s

Separate Nsight capture · n = 1

GENERATED VIDEO

960 × 544

124 frames · 24 FPS · 5.167 s

Reference configuration

PRO 5000 reference · 73,415 MiB / GPU

DiT TP1 / Ulysses8 · text TP8 · VAE tiles8

Local NCCL · RDMA disabled

Native VAE · FP16 decode autocast

Exact AdaLN table resident before timing

Fixed inputs / seed 43 · warmup excluded

CPU stage ranges · profiled request

Encoder	**7.272 s**
DiT	**204.814 s**
VAE / audio decode	**1.742 s**
Other / response	**1.140 s**

GPU work by operator

Accumulated CUDA durations, averaged over 8 GPUs

Operator	s / GPU	Share
NCCL (includes waiting)	**122.182**	**57.21%**
Attention	**47.105**	**22.06%**
GEMM	**39.650**	**18.57%**
Norm / activation / elementwise	2.288	1.07%
Layout / device copies	1.929	0.90%
CUDA memcpy / memset	0.064	0.03%
Convolution	0.192	0.09%
Other kernels	0.143	0.07%

NCCL includes synchronization and waiting.
Shares exclude CPU work / GPU idle gaps.

Host ranges and accumulated GPU work are different views; neither is a GPU critical-path decomposition.

NVIDIA

Five-Second Avatar: Live Conversation

Demo focus · reference identity, expression, native speech and continuity

01 · Natural presence

Listen with natural presence.
Microphone stays active; live captions show the question.

02 · Speech and expression

Generate a short answer as native audio and video.
Show generation time and actual reply-video start time.

03 · Continuous generation

Carry video and audio history across windows.
Inspect the face, mouth and transitions while speaking.

960 × 544 · 24 FPS · 8 steps · 8 SM120 GPUs · QAD NVFP4 / native FP4 VSA

119 new frames / 4.958 s per continuation · 20% history noise · video prefix off.
ASR and LLM share GPUs 6 / 7. The speech-to-video example is on slide 19.

NVIDIA

Three Ideas Behind Faster Generation

Technical overview · detailed algorithms and source excerpts remain available on hidden slides

Less model work

Few-step distillation reduces evaluations.
Sparse attention concentrates work on selected regions.
Quantized inference lowers arithmetic and memory cost.

Overlap across eight GPUs

Ulysses distributes attention across heads.
RDMA carries the exchanges.
Head bands overlap transfer with attention compute.

Keep the output path moving

CUDA Graph replays stable head-band compute.
Fused packing reduces intermediate copies.
NVFP4 VAE and MP4 pre-encoding reduce output cost.

These are optimization families. Use each demo’s own recipe and measurements; the four-step avatar and slide 18 are distinct workloads.

NVIDIA

Optimization Path: E2E Speedup

960 × 544 · 124 frames / 24 FPS · 8 exclusive GPUs · same inputs and seed 43

Stage	Cumulative change	Median (s)	E2E speedup vs. previous	RTF
00	Baseline · original BF16 / dense / 50 evaluations	213.402	—	41.304
01	+ PDD · 50 → 8 evaluations	44.327	4.814×†	8.579
02	+ Trained VSA · native 8-step model / runtime	25.915	1.710×	5.016
03	+ QAD NVFP4 DiT / native FP4 attention	19.460	1.332×	3.767
04	+ RDMA / direct layout / V-gate overlap	6.832	2.848×	1.322
05	+ 2 / 3 / 2 head pipeline	6.407	1.066×	1.240
06	+ Compute CUDA graphs	6.193	1.035×	1.199
07	+ 36-layer NVFP4 VAE / MP4 pre-encoding	5.420	1.143×	1.049
08–12	**+ Output and data-movement optimizations**	**4.922**	**1.101×**	**0.953**

08–12: fused tile / AAC preparation, direct I/O packing, compact routing and Q/K/V copy removal.

E2E speedup = previous displayed median ÷ current median. 08–12 compares the final result with stage 07.

Final 4.922 s · recheck 4.927 s · 6/6 requests < 5 s · 4.5 s target not reached

† Recorded median ratio; baseline refreshed separately (n = 2). Stages 01–12 retain n = 3 each; no isolated 50→8-step gain claimed.

RTF = latency / 5.167 s. Standalone H3; ASR / LLM excluded. Model, schedule and precision changes are not all lossless.

NVIDIA

Live Avatar: From a Question to a Video

960 × 544 · 24 FPS · 8 × SM120 GPUs · first window 5.167 s / each continuation adds 4.958 s

EXAMPLE · INTRODUCE YOURSELF

“Hi, could you introduce yourself?”

User question translated from the recorded demo

Microphone
Spoken question

→

ASR
Speech → text
Qwen3-ASR-0.6B

→

LLM
Text → short reply
Qwen3-4B

→

H3

→

Playback
Native audio + video
~5-second windows

LLM reply text

“Hi, I’m Miles. I help with simple, clear answers.”

H3 prompt · English excerpt

The man says exactly this English dialogue
in his reference voice:
"Hi, I'm Miles. I help with simple, clear answers."
… His lips follow the words.

A text reply becomes a speaking avatar clip.

Reference image



Input portrait supplied to H3

NVIDIA

Live Avatar: Recorded Conversation

74.55 seconds · user-recorded microphone conversation · native H3 audio + video



Play the recorded interaction · fullscreen avatar, live captions and timing overlays

NVIDIA


NVIDIA